

Advanced Machine Learning (Semester 2 2025)

Convolutional Neural Networks

Nina Hernitschek

Centro de Astronomía CITEVA
Universidad de Antofagasta

September 30, 2025

Workflow for Developing Neural Network Applications

1. understand and define the problem:
 - what are the inputs?
 - what are the expected outputs? (classification (binary, multiclass, multilabel); regression)
 - are inputs sufficiently informative to determine outputs?
2. define the data set
 - prepare data: normalization, PCA
3. choose a metric to measure success
4. Define evaluation methodology
 - e.g. cross-validation: holdout if plenty of data, otherwise K-fold
5. develop model
 - last-layer activation function
 - loss function
 - optimization
6. Train model to point of overfitting
7. Tune the model: regularization
8. Test the model
9. If sufficient: apply the model

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Recap

Software Frameworks



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Neural Network Architectures

A mostly complete chart of

Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org

Backfed Input Cell

Input Cell

Noisy Input Cell

Hidden Cell

Probabilistic Hidden Cell

Spiking Hidden Cell

Output Cell

Match Input Output Cell

Recurrent Cell

Memory Cell

Different Memory Cell

Kernel

Convolution or Pool

Perceptron (P)



Feed Forward (FF)




Radial Basis Network (RBF)




Deep Feed Forward (DFF)




Recurrent Neural Network (RNN)




Long / Short Term Memory (LSTM)



Gated Recurrent Unit (GRU)



Auto Encoder (AE)




Variational AE (VAE)



Denoising AE (DAE)



Sparse AE (SAE)



Markov Chain (MC)



Hopfield Network (HN)




Boltzmann Machine (BM)



Restricted BM (RBM)



Deep Belief Network (DBN)



Sasen Cain (@spectralradius)

Neural Network Architectures


'classic' trained feed-forward neural networks can handle object recognition



A subset of the MNIST Handwritten Digits dataset.

Neural Network Architectures


but we want not only object recognition...



Johnson et al., "Image Retrieval using Scene Graphs", CVPR 2015

Neural Network Architectures

but we want not only object recognition...



interacting galaxies
galaxy clusters
blending
gravitational lenses
foreground stars

...

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Neural Network Architectures

classic approach to **computer vision**:

treat image pixels as individual features to feed into the deep neural network (Multi-Layer Perceptron)



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Neural Network Architectures

classic approach to **computer vision**:

treat image pixels as individual features to feed into the deep neural network (Multi-Layer Perceptron)



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Neural Network Architectures

classic approach to **computer vision**:

treat image pixels as individual features to feed into the deep neural network (Multi-Layer Perceptron)



okay for simple and low-resolution images

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Neural Network Architectures

classic approach to **computer vision**:

treat image pixels as individual features to feed into the deep neural network (Multi-Layer Perceptron)



flattening the image removes spatial information
for complex and high-resolution dataset: Convolutional
Neural Network (CNN) is the solution

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Artificial Vision

Convolutional Neural Networks (CNN) were not invented overnight

topographical mapping of the visual cortex (in humans and animals):
nearby cells in the cortex represent nearby regions in the visual field



credit: Wandell et al. (2007)

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Artificial Vision

the visual cortex is organized hierarchically:



adapted from: Lane MJcIntosh, CS231n (2017)

simple cells: respond to light orientation

complex cells: respond to light orientation and movement

hypercomplex cells: respond to movement with an end point

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Artificial Vision

Neocognitron (Fukushima 1980)

uses a "sandwich" architecture of simple cells and complex cells where simple cells contain modifiable parameters, complex cells perform pooling operation



the first CNN

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolutional Neural Networks - The Idea

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolutional neural networks (CNNs) are similar to feedforward neural networks: A convolutional neural network consists of an input layer, hidden layers for feature learning and an output layer for classification.

The difference lies in the operations performed by the hidden layers: In a CNN, the hidden layers include layers that perform convolutions.

These networks harness principles from linear algebra, particularly matrix multiplication, to **identify patterns within an image**.

Convolutional Neural Networks - The Idea

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolutional neural networks (CNNs) are similar to feedforward neural networks: A convolutional neural network consists of an input layer, hidden layers for feature learning and an output layer for classification.

The difference lies in the operations performed by the hidden layers: In a CNN, the hidden layers include layers that perform convolutions.

These networks harness principles from linear algebra, particularly matrix multiplication, to **identify patterns within an image**.

Basic underlying idea:

An image is a matrix of pixels. Important parts on the image (objects, edges) are within close distances. Individual neurons respond to input only in a restricted region, like in the *visual cortex* of the brain.

Recap: Tensors

name	dimensionality (rank, order)	example
scalar	0	42
vector	1	(6.6 8 2)
matrix	2	$\begin{pmatrix} 1 & 2 & 3 \\ 55 & 0.6 & 1 \end{pmatrix}$
cube	3	$\begin{pmatrix} (6.7 \ 7.5 \ 2) & (4 \ 8.4 \ 5.1) & (55.7 \ 1 \ 2.5) \\ (7.5 \ 3 \ 2) & (3 \ 8 \ 9.1) & (6.5 \ 9.4 \ 2) \\ (2.7 \ 2 \ 2.5) & (1.4 \ 8 \ 1.5) & (6.2 \ 8.3 \ 2) \end{pmatrix}$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolutional Neural Networks - The Idea

A typical CNN architecture consists of alternating **Convolutional layers** and **Pooling layers** that extract features. Finally, their output is flattened into a one-dimensional feature layer that is passed into a Multi-Layer Perceptron (fully-connected layers) to obtain a prediction.



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
 fooling CNNs

Outlook

The Input



preserving the input image's spatial and temporal dependencies




In a CNN, the input is a tensor with shape
(input height) $\times$ (input width) $\times$ (input channels) $\times$ (number of inputs).

color image

movie

Convolution Layer



The objective of the Convolution Layer is to extract the high-level features, such as edges, from the input image.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



If you have used image/ photo editing, you likely have already applied convolutions, whether you realize it or not:

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook



blurring/smoothing




sharpening




edge detection




Convolution Layer




in feed-forward neural network:
a 32×32 pixel $\times$ 3 channels image is flattened



Convolution Layer



in a CNN:
processing a $32 \times 32 \text{ pixel} \times 3 \text{ channels}$ image as tensor



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook


Convolution Layer



in a CNN:
processing a 32×32 pixel $\times$ 3 channels image as tensor



Convolution Layer



What happens internally:

convolution is an element-wise multiplication of two matrices followed by summation

The operation performed in CNN is technically a **cross-correlation**, not a convolution:

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



What happens internally:

convolution is an element-wise multiplication of two matrices followed by summation

The operation performed in CNN is technically a **cross-correlation**, not a convolution:

Convolution:

$$y_{i,j} = \sum_l \sum_k x_{i-l, j-m} w_{l,m}$$

The diagram shows the input $x_{i-l, j-m}$ in an orange circle and the filter $w_{l,m}$ in a blue circle. Arrows point to them from the labels 'input' and 'filter' respectively.

the difference is for cross-correlation we don't have to 'flip' the filter's kernel relative to the input

Cross-Correlation (as carried out here):

$$y_{i,j} = \sum_l \sum_k x_{i+l, j+m} w_{l,m}$$

The diagram shows the input $x_{i+l, j+m}$ in an orange circle and the filter $w_{l,m}$ in a blue circle.

both can be interchanged through a simple rotation operation

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

convolved feature

4		

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1	1 _{x1}	1 _{x0}	0 _{x1}	0
0	1 _{x0}	1 _{x1}	1 _{x0}	0
0	0 _{x1}	1 _{x0}	1 _{x1}	1
0	0	1	1	0
0	1	1	0	0

convolved feature

4	3	

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$


image

1	1	1 _{x1}	0 _{x0}	0 _{x1}
0	1	1 _{x0}	1 _{x1}	0 _{x0}
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1	1	0
0	1	1	0	0

convolved feature

4	3	4

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

convolved feature

4	3	4
2		

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$


image

1	1	1	0	0
0	1 _{x1}	1 _{x0}	1 _{x1}	0
0	0 _{x0}	1 _{x1}	1 _{x0}	1
0	0 _{x1}	1 _{x0}	1 _{x1}	0
0	1	1	0	0

convolved feature

4	3	4
2	4	

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1	1	1	0	0
0	1	1 _{x1}	1 _{x0}	0 _{x1}
0	0	1 _{x0}	1 _{x1}	1 _{x0}
0	0	1 _{x1}	1 _{x0}	0 _{x1}
0	1	1	0	0

convolved feature

4	3	4
2	4	3

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

convolved feature

4	3	4
2	4	3
2		

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$


image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

convolved feature

4	3	4
2	4	3
2	3	

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image

1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}

convolved feature

4	3	4
2	4	3
2	3	4

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The process of convolution is basically the sum-product of the filter with an overlapping part of the image in a step-wise, strideful manner.

We illustrate this in the following **example**:

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ feature

kernel

$$K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

image


1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}

convolved feature

4	3	4
2	4	3
2	3	4

⇒ The kernel K shifts 9 times because of Stride Length = 1 (Non-Strided)

Convolution Layer



Design consideration: **The size of the convolution output.**

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Design consideration: **The size of the convolution output.**

The stride of the convolution (how many pixel the filter jumps each time) is not necessarily 1 pixel.

The size of the convolution output depends on implementation factors and might not be identical to the input.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Design consideration: **The size of the convolution output.**

Image size: $I \times I$

Filter size: $F \times F$

Stride: S

the **floor** operator: the smallest integer less than or equal to the input



Output size (in each dimension): $\lfloor (I - F) / S \rfloor + 1$

assuming the filter is not allowed to go beyond the edge

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Design consideration: **The size of the convolution output.**

Image size: $I \times I$

Filter size: $F \times F$

Stride: S


the **floor** operator: the smallest integer less than or equal to the input

Output size (in each dimension): $\lfloor (I - F) / S \rfloor + 1$
assuming the filter is not allowed to go beyond the edge



The result is a reduction in the output size, even for $S=1$
If this is not acceptable: Use **padding**.

Convolution Layer



The solution: **Zero Padding**

0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The solution: Zero Padding

0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

For an F width filter:

odd F : pad on left, right, top, bottom with a width of $(F - 1)/2$ zeros

even F : pad one side and top with $F/2$ zeros, and the other side and bottom with a width of $F/2 - 1$ zeros

The resulting image width is $I + F - 1$, the result of the convolution has width I

For a stride $S > 1$, zero padding is adjusted to ensure that the result of the convolution has width $\lceil I/S \rceil$: first zero-padding the image with $S\lceil I/S \rceil - I$ zeros and then apply the rules above.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
 fooling CNNs

Outlook

Convolution Layer



The **Computational Cost** of carrying out the convolution (cross-correlation) operation:

the cost of scanning an $I \times I$ image with a $F \times F$ filter is $\mathcal{O}(I^2 F^2)$:
we have F^2 multiplications at each of I^2 image pixels in the equation

$$y_{i,j} = \sum_l \sum_k x_{i+l,j+m} w_{l,m}$$

and carry this out $(I - F)/S$ times (the size of the output),
so the cost of scanning an $I \times I$ image with a $F \times F$ filter is $\mathcal{O}(I^2 F^2)$

despite also $(I - F)/S$ (how often we carry out this operation as the filter with size F moves with stride S over the image with size I) contributes to the computing time, for the $\mathcal{O}$ we only give the highest order

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The **Computational Cost** of carrying out the convolution (cross-correlation) operation:

the cost of scanning an $I \times I$ image with a $F \times F$ filter is $\mathcal{O}(I^2 F^2)$:
we have F^2 multiplications at each of I^2 image pixels in the equation

$$y_{i,j} = \sum_l \sum_k x_{i+l,j+m} w_{l,m}$$

and carry this out $(I - F)/S$ times (the size of the output),
so the cost of scanning an $I \times I$ image with a $F \times F$ filter is $\mathcal{O}(I^2 F^2)$



expensive for large filters

a solution: convolutions using **Discrete Fourier Transforms (DFTs)**:

$$Y = \text{IDFT2}(\text{DFT2}(X) \circ \text{conj}(\text{DFT2}(W)))$$

where IDFT2 is the **second order inverted discrete Fourier transform**
computational cost $\mathcal{O}(I^2 \log F)$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
 fooling CNNs

Outlook

Convolution Layer



The **Computational Cost** of carrying out the convolution (cross-correlation) operation:

we have F^2 multiplications at each of I^2 image pixels in the equation

$$y_{i,j} = \sum_l \sum_k x_{i+l,j+m} w_{l,m}$$

and carry this out $(I - F)/S$ times (the size of the output), so the cost of scanning an $I \times I$ image with a $F \times F$ filter is $\mathcal{O}(I^2 F^2)$



expensive for large filters

a solution: convolutions using **Discrete Fourier Transforms (DFTs)**:

$$Y = IDFT2(DFT2(X) \circ conj(DFT2(W)))$$

where IDFT2 is the **second order inverted discrete Fourier transform** computational cost $\mathcal{O}(I^2 \log F)$



significant reduction in computational cost

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



How to **set** the kernel values of the filters?

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



How to **set** the kernel values of the filters?

In the Convolutional Neural Network, these kernel values in the filter are **essentially weights to be learned and trained** to automatically capture spatial edge patterns in images.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



How to **set** the kernel values of the filters?

In the Convolutional Neural Network, these kernel values in the filter are **essentially weights to be learned and trained** to automatically capture spatial edge patterns in images.

In addition to its ability to capture complex patterns, the small number of weights (filters are usually small relative to the images) in each layer makes the CNN very efficient to train.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook


Convolution Layer



After each convolution operation, a **Rectified Linear Unit (ReLU) transformation** is applied to the feature map, introducing nonlinearity to the model.

This is the activation function we have seen before.


ReLU effectively removes negative values from an activation map by setting them to zero.




$$\varphi(z) = \max(0, z)$$

$$\varphi'(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{otherwise.} \end{cases}$$

Convolution Layer



A hypothetical, simple Convolutional Layer:



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



A hypothetical, simple Convolutional Layer:



The number of channels in a Convolutional Layer output equals the number of 3D (length $\times$ width $\times$ channels) filters used in the layer.

During the 3D convolutions, each output channel, as a 2D matrix, has additional bias broadcasted to each element, before a ReLU activation function is applied.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Three **hyperparameters** control the size of the output volume of the convolutional layer:

1. The **filter depth** affects the depth of the output. E.g.: three distinct filters would yield three different feature maps.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Three **hyperparameters** control the size of the output volume of the convolutional layer:

1. The **filter depth** affects the depth of the output. E.g.: three distinct filters would yield three different feature maps.
2. **Stride** is the number of pixels that the kernel moves over the input matrix. A larger stride means smaller overlap of receptive fields and smaller spatial dimensions of the output volume.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



Three **hyperparameters** control the size of the output volume of the convolutional layer:

1. The **filter depth** affects the depth of the output. E.g.: three distinct filters would yield three different feature maps.
2. **Stride** is the number of pixels that the kernel moves over the input matrix. A larger stride means smaller overlap of receptive fields and smaller spatial dimensions of the output volume.
3. **Zero-padding** is usually used when the filters do not fit the input image. This sets all elements that fall outside of the input matrix to zero.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The total **number of parameters to be trained**:

Recap

Motivation

Convolutional
Neural
Networks

**Convolution
Layer**

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The total **number of parameters to be trained**:

Considering if there are 10 filters that are $3 \times 3 \times 3$ in dimension, then the total number of parameters (weights and biases) in the Convolutional Layer would be: Total Parameters = $(3 \times 3 \times 3 + 1) \times 10 = 280$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The total **number of parameters to be trained**:

Considering if there are 10 filters that are $3 \times 3 \times 3$ in dimension, then the total number of parameters (weights and biases) in the Convolutional Layer would be: Total Parameters = $(3 \times 3 \times 3 + 1) \times 10 = 280$

The total number of parameters usually increase in later Convolutional Layers, along with the typical increase in channel size and the number of filters.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Convolution Layer



The total **number of parameters to be trained**:

Considering if there are 10 filters that are $3 \times 3 \times 3$ in dimension, then the total number of parameters (weights and biases) in the Convolutional Layer would be: Total Parameters = $(3 \times 3 \times 3 + 1) \times 10 = 280$

The total number of parameters usually increase in later Convolutional Layers, along with the typical increase in channel size and the number of filters.

computing the total number of parameters in a layer:

I input size

F filter size

P padding size

S stride size

n number of filters

n_{prev} number of filters in previous layer

output dimensions: $\lfloor \frac{I+2P-F}{S} + 1 \rfloor^2$

filter dimension: $F^2 \times n_{prev}$

number of weights: $F^2 \times n_{prev} \times n$

number of biases n

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Pooling Layer



Each Convolutional Layer is followed by a Pooling Layer. Similar to the Convolution Layer, the Pooling Layer reduces the spatial size of the Convolved Feature. It suppresses noise and extracts dominant features.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Pooling Layer



Each Convolutional Layer is followed by a Pooling Layer. Similar to the Convolution Layer, the Pooling Layer reduces the spatial size of the Convolved Feature. It suppresses noise and extracts dominant features.

In detail, a pooling layer carries out the following tasks:

- Extract the most important (relevant) features by getting the maximum number or averaging the numbers.
- Reduce the dimensionality (number of pixels) of the output returned from previous convolutional layers.
- Reduce the number of parameters in the network.
- Remove any noise present in the features extracted by previous convolutional layers.
- Increase the accuracy of CNNs.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Pooling CNNs

Outlook


Pooling Layer



The pooling layer contains three elements: Feature Map, Filter and Pooled Feature Map.

The **Pooling Operation** happens between a section of the feature map and the filter. It outputs the pooled feature map which is a downsampled version of the feature map.

Example: Maximum Pooling here: with a 2×2 filter and $S = 2$



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
 fooling CNNs

Outlook


Pooling Layer




alternatives:

p-Norm Pooling

here: with a 2×2 filter and $S = 2$, $p = 5$




Pooling Layer



alternatives:


p-Norm Pooling

here: with a 2×2 filter and $S = 2$, $p = 5$



Average Pooling

here: with a 2×2 filter and $S = 2$



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Pooling CNNs

Outlook

Pooling Layer



Other than convolutional filters which are learned, pooling filters are typically not learned.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Pooling Layer



Other than convolutional filters which are learned, pooling filters are typically not learned.

But it is still possible to do so!

In this case, the same network is applied on each block when moving the pooling filter.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Pooling Layer



Which pooling should we chose?

The most common pooling layer is **Max Pooling**, with a pooling window size of 2 and stride of 2 - resulting in a output dimension that is halved. By picking the most activated value in the pooling window, the Max Pooling layer removes redundant features and helps in preventing over-fitting.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Pooling Layer



Which pooling should we chose?

The most common pooling layer is **Max Pooling**, with a pooling window size of 2 and stride of 2 - resulting in a output dimension that is halved. By picking the most activated value in the pooling window, the Max Pooling layer removes redundant features and helps in preventing over-fitting.

On the other hand, **Average Pooling** typically happens between the Convolutional Layers and the Fully Connected Layers, in place of the Flatten operation. Average Pooling is sometimes preferred compared to the Flatten operation due to it having less features passed into the Fully Connected Layers, and may reduce over-fitting.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Pooling CNNs

Outlook

Fully Connected Layers



The fully connected layers classify the detected features in the image into a class label. These are the final layers in a CNN.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Fully Connected Layers



The fully connected layers classify the detected features in the image into a class label. These are the final layers in a CNN.

As input, the previous layer is flattened. There can be multiple fully connected layers. The final layer carries out the classification.

The ReLU activation is used in each fully connected layer except in the final layer in which we use the Softmax activation for multiclass classification.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook


Fully Connected Layers



A **Flatten Operation** is necessary:

In a CNN, the final pooled feature map is fed to a Multilayer Perceptron (MLP) to classify the feature map it into a class label.

MPLs only accept one-dimensional data - this requires flattening the final pooled feature map into a single column. Unlike flattening the original image, important pixel dependencies are retained when pooled maps are flattened.



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training


CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook


Fully Connected Layers



The Fully Connected Layers learn non-linear combinations of the high-level features as represented by the output of the convolutional layer.

Neurons in a fully connected layer have connections to all activations in the previous layer, as seen in regular (non-convolutional) artificial neural networks.

The **activation function** is usually a softmax function:



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary: The Building Blocks of a CNN

The CNN is built along how the vision is processed in mammals.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary: The Building Blocks of a CNN

The CNN is built along how the vision is processed in mammals.

It includes:

- convolutional layers comprising learned filters that scan the outputs of previous layers
- pooling layers that downsample outputs from convolutional layers
- an output layer that works on flattened data from the last pooling layer.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary: The Building Blocks of a CNN

The CNN is built along how the vision is processed in mammals.

It includes:

- convolutional layers comprising learned filters that scan the outputs of previous layers
- pooling layers that downsample outputs from convolutional layers
- an output layer that works on flattened data from the last pooling layer.

Convolution can change the size of the output. This can be controlled with zero padding.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary: The Building Blocks of a CNN

The CNN is built along how the vision is processed in mammals.

It includes:

- convolutional layers comprising learned filters that scan the outputs of previous layers
- pooling layers that downsample outputs from convolutional layers
- an output layer that works on flattened data from the last pooling layer.

Convolution can change the size of the output. This can be controlled with zero padding.

Pooling layers may perform operations such as max, p-norm, or learned downsampling networks.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

As in the case of regular MLP, training happens through backpropagation. The only difference comes from the structure of the neural network.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

As in the case of regular MLP, training happens through backpropagation. The only difference comes from the structure of the neural network.

As in the case of a regular MLP:

- a training sample is provided
- the difference between the desired output and the true output is defined and calculated in response to any output of the training sample
- network parameters are trained through backpropagation with gradient descent

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

CNNs give us two key benefits: local invariance and compositionality.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

CNNs give us two key benefits: local invariance and compositionality.

In the final **multi-layer perceptron**, all weights and biases for classification are learned to make predictions based on detected higher-level features.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

CNNs give us two key benefits: local invariance and compositionality.

In the final **multi-layer perceptron**, all weights and biases for classification are learned to make predictions based on detected higher-level features.

In the **convolutional layers**, the filters must learned. In the context of image classification, our CNN may learn to:

- Detect edges from raw pixel data in the first layer.
- Use these edges to detect shapes (i.e., *blobs*) in the second layer.
- Use these shapes to detect higher-level features such as car parts, lensed galaxies etc.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

CNNs give us two key benefits: local invariance and compositionality.

In the final **multi-layer perceptron**, all weights and biases for classification are learned to make predictions based on detected higher-level features.

In the **convolutional layers**, the filters must learned. In the context of image classification, our CNN may learn to:

- Detect edges from raw pixel data in the first layer.
- Use these edges to detect shapes (i.e., *blobs*) in the second layer.
- Use these shapes to detect higher-level features such as car parts, lensed galaxies etc.

The **number of parameters**:

Let each layer j have K_j maps.

Let the filters in the j th layer be of size $L_j \times L_j$

For the j th layer we will require $K_j(K_{j-1}L_j^2 + 1)$ filter parameters.

The total number of parameters for each convolutional layer is then

$$\sum_j K_j(K_{j-1}L_j^2 + 1)$$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

Computing the Loss

Given a training set with input-output pairs:

$$(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots, (x_T, \hat{y}_T)$$

The loss on the i th instance is $\Delta(x_i, \hat{y}_i)$.

The total loss is $E = \frac{1}{T} \sum_{i=1}^T \Delta(x_i, \hat{y}_i)$.

The loss is minimized w.r.t. the weights and biases θ

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Training CNN

Computing the Loss

Given a training set with input-output pairs:

$$(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots, (x_T, \hat{y}_T)$$

The loss on the i th instance is $\Delta(x_i, \hat{y}_i)$.

The total loss is $E = \frac{1}{T} \sum_{i=1}^T \Delta(x_i, \hat{y}_i)$.

The loss is minimized w.r.t. the weights and biases θ

Training through gradient descent

initialize all weights and biases θ

for every layer l for all filter indices m , until loss has converged update

$\theta_k \rightarrow \theta_{k+1}$:

$$\theta_{k+1}(l, m, j, x, y) = \theta_k(l, m, j, x, y) - \eta \frac{\partial E}{\partial \theta_k(l, m, j, x, y)}$$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs


Outlook

Training CNN

Backpropagation continues in the usual manner until the computation of the derivative of the divergence w.r.t. its inputs to the first "flat" layer

Backpropagation from the flat MLP requires special consideration:

- the shared computation in the convolution layers
- the pooling layers



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Pooling CNNs

Outlook

Training CNN

Both the Forward pass and the Backpropagation of a Convolutional layer are Convolutions.

So the gradients for backpropagation become, with filter F , layer output y_l and layer input $x = y_{l-1}$:

$$\frac{\partial E}{\partial F} = \text{Conv} \left(\text{input } x, \text{loss gradient } \frac{\partial E}{\partial y} \right)$$

$$\frac{\partial E}{\partial x} = \text{Conv} \left(180^\circ \text{ rotated filter } F', \text{loss gradient } \frac{\partial E}{\partial y} \right)$$

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary

CNN include convolutional layers comprising learned filters that scan the outputs of the previous layer.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary

CNN include convolutional layers comprising learned filters that scan the outputs of the previous layer.

Pooling layers operate on groups of outputs from the convolutional layer to reduce the network size.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

CNN include convolutional layers comprising learned filters that scan the outputs of the previous layer.

Pooling layers operate on groups of outputs from the convolutional layer to reduce the network size.

The parameters of the CNN can be learned through backpropagation.

Applications of Convolutional Neural Networks

Advantages of Convolution Neural Networks:

- used for deep learning with few parameters
- less parameters to learn as compared to fully connected layer
- tool for image processing

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Advantages of Convolution Neural Networks:

- used for deep learning with few parameters
- less parameters to learn as compared to fully connected layer
- tool for image processing

Disadvantages of Convolution Neural Networks:

- Comparatively complex to design and maintain
- Comparatively slow [depends on the number of hidden layers]

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Advantages of Convolution Neural Networks:

- used for deep learning with few parameters
- less parameters to learn as compared to fully connected layer
- tool for image processing

Disadvantages of Convolution Neural Networks:

- Comparatively complex to design and maintain
- Comparatively slow [depends on the number of hidden layers]

General Applications of Convolution Neural Networks:

- Image processing
- Computer Vision
- Speech Recognition

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Specific Applications in Astronomy:

1. extracting prominent features from astronomical images, e.g.
 - craters
 - gravitational lenses

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Specific Applications in Astronomy:

1. extracting prominent features from astronomical images, e.g.
 - craters
 - gravitational lenses
2. classifying objects on astronomical images

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Specific Applications in Astronomy:

1. extracting prominent features from astronomical images, e.g.
 - craters
 - gravitational lenses
2. classifying objects on astronomical images
3. classifying time-series data based on a plot (light curve)

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Applications of Convolutional Neural Networks

Specific Applications in Astronomy:

example:



left: the Hubble Heritage image of the Carina Nebula

right: the locations of each of the 2240 sub images in the clustering space

Convolutional Neural Networks in Astronomy, and Applications for Diffuse Structure Discovery, Peek, J. E. G.; Jones, Craig K.; Hargis, Jonathan, *Astronomical Data Analysis Software and Systems XXVII*, April 2020

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Astronomy with CNN

Rotation-invariant convolutional neural networks for galaxy morphology prediction, S. Dieleman, K. W. Willett, J. Dambre (2015)

Measuring the morphological parameters of galaxies is a key requirement for studying their formation and evolution. The deep neural network model for galaxy morphology classification exploits translational and rotational symmetry.





Figure 4: Schematic overview of the CNN architecture. The input (1) is rotated/flipped for different viewpoints (2), which are cropped to reduce redundancy (3). Processing is done by convolutional and pooling layers (4), their output is concatenated and processed by dense layers (5) to obtain predictions (6).


Astronomy with CNN

CMU DeepLens: Deep Learning For Automatic Image-based Galaxy-Galaxy Strong Lens Finding, F. Lanusse, Qu. Ma, N. Li, et al. (2017)

CMU DeepLens is a fully automated galaxy-galaxy lens finding method based on Deep Learning, trained and validated on 20,000 LSST-like mock observations including a range of lensed systems of various sizes and signal-to-noise ratios.



(a) Single best epoch images



(b) Stack images

Figure 7. S/N and Einstein radius distribution of the simulated and recovered lens populations, for our fiducial 1% FPR detection threshold.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Astronomy with CNN

Searching for exoplanets using artificial intelligence, K. A. Pearson, L. Palafox, C. A. Griffith (2018)

the ideal algorithm should be

- fast
- robust to noise
- capable of learning and abstracting highly nonlinear systems





Figure 2. Past transit detection algorithms are accomplished by correlating the data with a simple box model through a least-squares optimization.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Astronomy with CNN

Finding Strong Gravitational Lenses in the Kilo Degree Survey with Convolutional Neural Networks, C. E. Petrillo, C. Tortora, S. Chatterjee et al. (2017)

A morphological classification method based on a Convolutional Neural Network (CNN) for recognizing strong gravitational lenses is applied to the 255 square degrees of the Kilo Degree Survey (KiDS), one of the current-generation optical wide surveys.




Figure 3. A schematic of the training-set creation.

Astronomy with CNN

Quasar microlensing light curve analysis using deep machine learning, G. Vernardos & G. Tsagkatakis (2019)

Objective: Use CNN to model quasar microlensing light curves and extract the size and temperature profile of the accretion disc.

The effects of microlensing are simulated by magnification maps: pixellated representations of the source plane caustics due to the microlenses.




Figure 2. Mock light curve from magnification map. The light curve trajectory on the map is depicted as a red arrow.

CNN Architectures

How to construct a CNN?

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

**CNNs in
Astronomy**

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

CNN Architectures

How to construct a CNN?



There are several common CNN architectures we can adapt to solve our science questions.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


**CNN
Architectures**

Visualizing
Features &
Fooling CNNs

Outlook

Transfer Learning for Images

The **basic concept** of transfer learning is training on a large dataset and transfer the knowledge to a smaller dataset. It is based on the idea that convolutional layers extract general features applicable across images.



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Transfer Learning for Images

The **basic concept** of transfer learning is training on a large dataset and transfer the knowledge to a smaller dataset. It is based on the idea that convolutional layers extract general features applicable across images.




Following is the general outline for transfer learning for object recognition:

1. Load in a pre-trained CNN model trained on a large dataset.
2. Freeze parameters (weights) in model's lower convolutional layers.
3. Add custom classifier with several layers of trainable parameters.
4. Train classifier layers on training data available for task.
5. Fine-tune hyperparameters and unfreeze more layers as needed.

LeNet-5

LeNet-5 was proposed by Y. LeCun (1998) in **Gradient-Based Learning Applied to Document Recognition**. This CNN is used to predict handwritten numeric characters in greyscale images.



Excluding pooling, LeNet-5 consists of 5 layers:

- 2 convolution layers with kernel size 5×5 , followed by
- 3 fully connected layers.

Tanh activation function is used except for the last layer (which has softmax). LeNet-5 has 60,000 trainable parameters.

Input: 32×32 pixel image. Largest character is 20×20 . Pixel values are normalized (mean of pixels = 0, std of pixels = 1).



GoogLeNet (Inception-v1)

GoogLeNet, also called Inception-v1 (Szegedy et al. 2014) is built on the concept of a **inception module**: design a good local network topology and then **stack** these modules on top of each other (network within a network)

a naive inception module:


apply parallel filter operations on the input from previous layer:

- multiple receptive field sizes for convolution (1×1 , 3×3 , 5×5)
- pooling operation (3×3)



GoogLeNet (Inception-v1)

the problem: **computational complexity**



the **convolutional operations**:

e.g. with a module input of $28 \times 28 \times 256$, we get:


1x1 convolution: $28 \times 28 \times 128 \times 1 \times 1 \times 256$ operations

3x3 convolution: $28 \times 28 \times 192 \times 3 \times 3 \times 256$ operations

5x5 convolution: $28 \times 28 \times 96 \times 5 \times 5 \times 256$ operations

GoogLeNet (Inception-v1)

the problem: **computational complexity**



the **convolutional operations**:

e.g. with a module input of $28 \times 28 \times 256$, we get:

1x1 convolution: $28 \times 28 \times 128 \times 1 \times 1 \times 256$ operations

3x3 convolution: $28 \times 28 \times 192 \times 3 \times 3 \times 256$ operations

5x5 convolution: $28 \times 28 \times 96 \times 5 \times 5 \times 256$ operations

$\Rightarrow$ total 854×10^6 operations

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

GoogLeNet (Inception-v1)

the problem: **computational complexity**



the **convolutional operations**:

e.g. with a module input of $28 \times 28 \times 256$, we get:

1x1 convolution: $28 \times 28 \times 128 \times 1 \times 1 \times 256$ operations

3x3 convolution: $28 \times 28 \times 192 \times 3 \times 3 \times 256$ operations

5x5 convolution: $28 \times 28 \times 96 \times 5 \times 5 \times 256$ operations

⇒ total 854×10^6 operations

plus: **pooling layer** preserves the feature depth which means total depth after concatenation can only grow at every layer

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook


GoogLeNet (Inception-v1)

Solution (Szegedy et al. 2014): **bottleneck layers** that use 1×1 convolutions to reduce feature depth



naive

vs.



with dimensionality reduction

computationally more efficient:

e.g. with a module input of $28 \times 28 \times 256$, we get:

total 358×10^6 operations (compared to 854×10^6 operations for the naive version)

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

GoogLeNet (Inception-v1)

adding more layers comes with a **drawback**: exploding/vanishing gradients

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer


Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook



GoogLeNet (Inception-v1)

adding more layers comes with a **drawback**: exploding/vanishing gradients
GoogLeNet solves this by adding **auxiliary classifiers** after the 3rd and 6th inception module to increase the gradient signal that gets propagated back. During training, their loss is added to the network's total loss with a 0.3 weight. At inference time, these auxiliary networks are discarded.




GoogLeNet (Inception-v1)

adding more layers comes with a **drawback**: exploding/vanishing gradients



GoogLeNet solves this by adding **auxiliary classifiers** after the 3rd and 6th inception module to increase the gradient signal that gets propagated back. During training, their loss is added to the network's total loss with a 0.3 weight. At inference time, these auxiliary networks are discarded. The auxiliary classifiers, as well as the last layer of the main classifier, use a softmax activation function (all others use ReLU).

Auxiliary classifiers start with a 5×5 average-pooling, followed by a convolution layer with 1×1 kernel size and two fully connected layers.



CNN Architectures - Case Studies

comparing complexity



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

CNN Architectures - Summary


CNN Architecture	Default Input	Default Output	Number of Layers	Number of Parameters	Activation Function	specific properties
LeNet-5	$32 \times 32 \times 1$	10	5	60k	tanh	convolution layer
AlexNet	$244 \times 244 \times 3$	1000	8	60M	ReLU	local response normalization
VGG-16	$244 \times 244 \times 3$	1000	16	138M	ReLU	very deep but single-thread
GoogLeNet	$244 \times 244 \times 3$	1000	22	7M	ReLU	inception module, auxiliary classifiers
ResNet-50	$244 \times 244 \times 3$	1000	50	26M	ReLU	batch normalization, residual blocks

local response normalization:

This implements the idea of lateral inhibition where an excited neuron inhibits its neighbours, leading to contrast in that area.

residual blocks:

Traditionally, each layer feeds into the next layer. In a network with residual blocks, each layer feeds into the next layer and into the layers about 2-3 hops away.



Visualizing CNN Features

Visualizing how a CNN learns to identify different features present in images provides a deeper insight into the model.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer


Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs


Outlook



Visualizing CNN Features

Visualizing how a CNN learns to identify different features present in images provides a deeper insight into the model.

This can be done with a process called **Gradient Ascent**.




Visualizing CNN Features

Visualizing how a CNN learns to identify different features present in images provides a deeper insight into the model.

This can be done with a process called **Gradient Ascent**.

Gradient Ascent starts with a zero or random image and continuously backpropagates until an image is generated that maximizes the score for the particular class (e.g.: dog).



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Fooling a CNN

a CNN's classification:

African elephant



koala



schooner



iPod



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures


Visualizing
Features &
Fooling CNNs

Outlook

Fooling a CNN

Suppose we feed in an image to a CNN. Instead of maximizing the likelihood of the correct class (e.g. elephant) we maximize an incorrect class (e.g. koala) while making minimal changes to the image.

African elephant




koala



Difference



10x Difference




schooner




iPod



Difference



10x Difference



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Fooling a CNN

Suppose we feed in an image to a CNN. Instead of maximizing the likelihood of the correct class (e.g. elephant) we maximize an incorrect class (e.g. koala) while making minimal changes to the image.



Intriguing properties of neural networks, Szegedy et al. (2013)
Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images, Nguyen, Yosinski, Clune (2014)

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy


CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Fooling a CNN

Suppose we feed in an image to a CNN. Instead of maximizing the likelihood of the correct class (e.g. elephant) we maximize an incorrect class (e.g. koala) while making minimal changes to the image.



Intriguing properties of neural networks, Szegedy et al. (2013)
Neural Networks are Easily Fooled: High Confidence Predictions for
Unrecognizable Images, Nguyen, Yosinski, Clune (2014)



Notice that the changes are so minimal that the two images are indistinguishable to humans


Fooling a CNN

Why does this work?

Using the training images we fit some **decision boundaries** in the high-dimensional parameter space.

While doing so based on the **finite set of training images**, also decisions about many unseen points in parameter space are made.

As a result: Of the many points in this high-dimensional parameter space belonging to a certain class, only a few represent true images.



credit:
Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images, Nguyen, Yosinski, Clune (2014)

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Google's *Deep Dream*

A fun way to use CNN:



Exaggerating feature attributes or textures using information that the `bvlc_googlenet` model learned during training.



The Mona Lisa with DeepDream effect using VGG16 network trained on ImageNet.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Google's *Deep Dream*

A fun way to use CNN:

1. Forward propagate a baseline image through your model. Compute the activations at a chosen layer.
2. To maximize the neuron activations in the layer we care about, set the gradient of chosen layer equal to its activation.
3. During the backward pass, compute the gradient on the loss w.r.t. the input image, as in the case of gradient ascent.
4. Execute the update rule to update your input image. We change the image so that these neurons fire even more.



you can try it out at: <https://deepai.org/>



Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary and Outlook

CNNs enable us to deal with images more naturally.

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

Summary and Outlook

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook

CNNs enable us to deal with images more naturally.

Next time we look at another neural network architecture:
recurrent neural networks (RNN).

Summary and Outlook

October 21: **Presentation on project idea**

Presentation (15 minutes) should cover:

1. Clear explanation of the overall problem you want to solve
2. The dataset(s) you will use
2. Relationship to the topics covered in class: Which models/algorithms are you planning to use, how will you evaluate performance
4. List of papers you plan to read as references
5. How will you structure the project, rough timeline

Recap

Motivation

Convolutional
Neural
Networks

Convolution
Layer

Pooling Layer

Training

CNNs in
Astronomy

CNN
Architectures

Visualizing
Features &
Fooling CNNs

Outlook